

Applied Math with MATLAB

CME 192 LECTURE 3

01/22/2026

Outline

Numerical Linear Algebra

Sparse matrices

Matrix decomposition

Linear system solvers

ODE and PDE

Classification of ODEs, PDEs

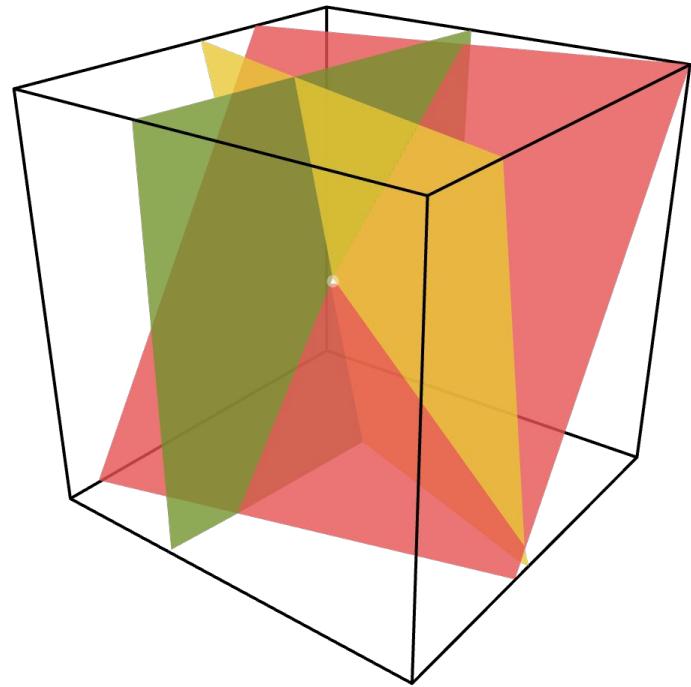
Numerical Methods for Solving

Geometry Definitions + Workflow

Symbolic Math

Numerical Linear Algebra

Applied Math with MATLAB



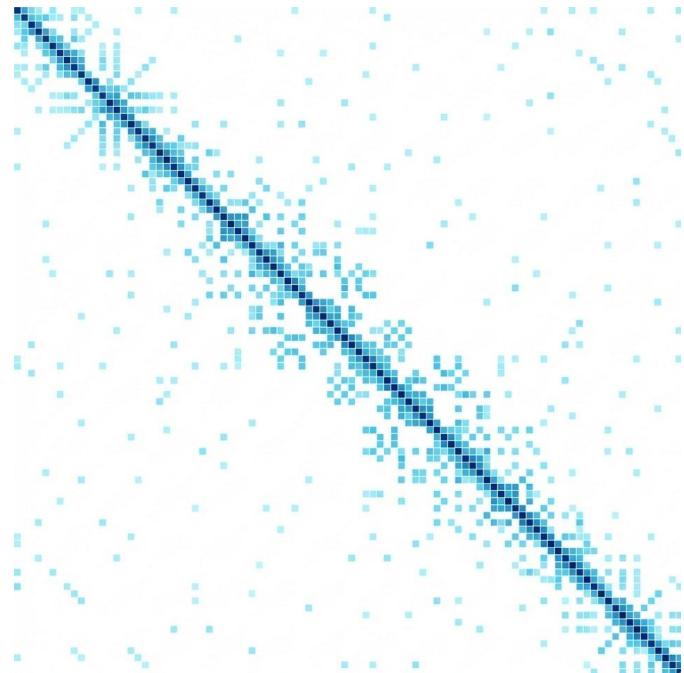
Dense vs. Sparse Matrices

This is a **sparse matrix**, a matrix with relatively small number of nonzero entries, compared to its size.

Let $A \in R^{m \times n}$ be a sparse matrix with v nonzeros.

Dense storage requires mn entries.

Sparse format saves storage.



Dense Matrix

1.2	0.5	-0.1	1.1	1.8	-4.3
3.7	-1.1	5.0	2.6	-3.4	4.8
1.6	-2.2	0.3	5.5	1.9	-0.9
-0.8	1.2	3.2	0.5	-4.3	1.7
1.9	-0.8	3.2	-5.1	2.4	6.5
1.3	-4.7	0.2	-3.8	4.1	3.8

Sparse Matrix

3.7	0	0	0	0	0
0	0	5.3	0	0	0
0	4.5	0	0	0	0
0	0	0	0	0	0
0	0	0	0	1.4	0
0	0	0	0	0	0

Sparse Matrix Storage Formats

1. Triplet format

Store nonzero values and corresponding row/column

Storage required = $3v$ ($2v$ ints and v doubles)

General in that no assumptions are made about sparsity structure

Accepted as input for sparse matrix constructions by MATLAB

The diagram illustrates the conversion of a sparse matrix into its triplet format. On the left, a 5x6 matrix is shown with non-zero elements highlighted in blue. An arrow points from this matrix to a table on the right, which represents the triplet format.

Row	0	0	1	1	3	3
Column	2	4	2	3	1	2
Value	3	4	5	7	2	6

Example: Dense vs. sparse storage (tridiagonal 1000×1000)

Dense (double): $1000 \times 1000 = 1,000,000$ entries

$\approx 8,000,000$ bytes (assuming 8 bytes per double)

Triplet storage: $v = 3n - 2 = 2998 \rightarrow 3v = 8994$ numbers

$\approx 71,952$ bytes if each stored number is 8 bytes

Note: actual sparse storage depends on data types (indices vs values).

```
A = zeros(n,n);
A(1:n+1:end) = 2;
A(2:n+1:end) = -1;
A(n+1:n+1:end) = -1;

I = [ (1:n)';      (1:n-1)';     (2:n)'    ];
J = [ (1:n)';      (2:n)';      (1:n-1)'  ];
V = [ 2*ones(n,1); -ones(n-1,1);
      -ones(n-1,1) ];

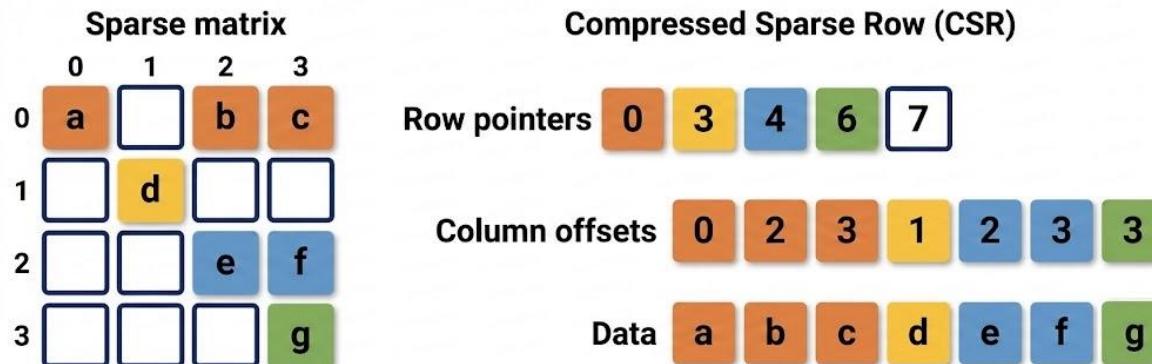
S = sparse(I,J,V,n,n);
whos A S I J V
```

Name	Size	Bytes	Class	Attributes
A	1000x1000	8000000	double	
I	2998x1	23984	double	
J	2998x1	23984	double	
S	1000x1000	55976	double	sparse
V	2998x1	23984	double	

2. Compressed Sparse Row (CSR) format

Store nonzero values, corresponding column, and pointer into value array corresponding to first nonzero in each row

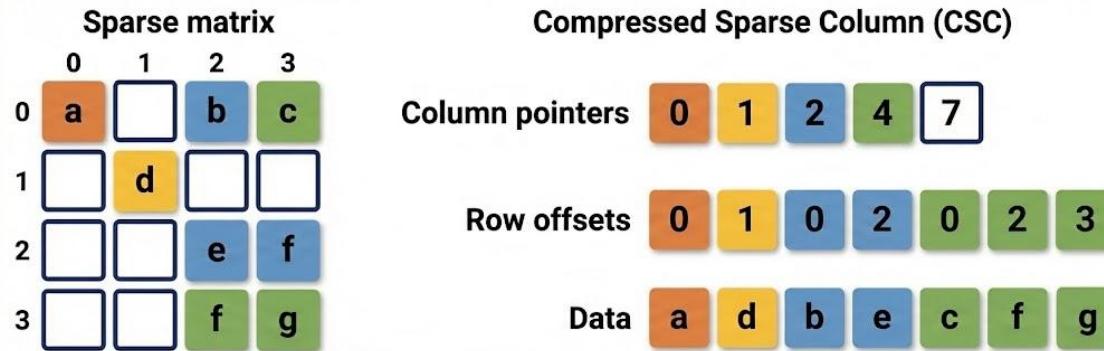
$$\text{Storage required} = 2v + m$$



3. Compressed Sparse Column (CSC) format

Store nonzero values, corresponding row, and pointer into value array corresponding to first nonzero in each column

Storage required = $2v + n$



4. Others

Diagonal Storage format (useful for banded matrices)

Skyline Storage format

Block Compressed Sparse Row (BSR) format

$$\begin{pmatrix} 1 & 2 & 0 & 0 & 0 & 0 \\ 3 & 4 & 5 & 0 & 0 & 0 \\ 0 & 6 & 7 & 0 & 0 & 0 \\ 8 & 0 & 9 & 10 & 11 & 0 \\ 0 & 13 & 0 & 0 & 14 & 15 \\ 0 & 0 & 16 & 0 & 17 & 18 \end{pmatrix}$$

(a) A sparse matrix

VAL			
-	-	1	2
-	3	4	5
-	6	7	0
8	9	10	11
13	0	14	15
16	17	18	-

OFFSET

-3	-1	0	1
----	----	---	---

(b) Storage in diagonal format

1	0	6	7	*	*
2	1	8	2	*	*
*	*	1	4	*	*
*	*	5	1	*	*
*	*	4	3	7	2
*	*	0	0	0	0

Diagonal Storage Format

BSR Format

Break-Even Point for Sparse Storage

For $A \in R^{m \times n}$ with v nonzeros, there is a value of v where sparse vs dense storage is more efficient.

- For the triplet format, the cross-over point is defined by

$$3v = mn,$$

- If $v \leq mn/3$ use sparse storage, otherwise use dense format.
- Cross-over point depends not only on m, n, v , but also on the data types of row, col, val.
- Besides storage efficiency, data access for linear algebra applications and ability to exploit symmetry in storage is also important.

Quiz

Suppose you have a tridiagonal 1000x1000 matrix. All entries in those three diagonals are non-zero. Calculate how many numbers can we omit if we store it as a triplet instead of a full matrix?

It sounds that a tridiagonal matrix is better stored as a triplet. For a tridiagonal $n \times n$ matrix, find the greatest n in which the triplet is not helpful

in saving memory for a tridiagonal matrix where all entries in those diagonals are non-zero.

Quiz

Suppose you have a tridiagonal 1000x1000 matrix. All entries in those three diagonals are non-zero. Calculate how many numbers can we omit if we store it as a triplet instead of a full matrix?

It sounds that a tridiagonal matrix is better stored as a triplet. For a tridiagonal $n \times n$ matrix, find the greatest n in which the triplet is not helpful in saving memory for a tridiagonal matrix where all entries in those diagonals are non-zero.

Solution

full matrix: $1000 \times 1000 = 10^6$

triplet: $v = 2998$, $3v = 8994$

For a tridiagonal matrix, $v = 3n - 2$. Solving $3v = n^2$ gives $n \approx 8.275$.

Timing example (MATLAB): dense vs. sparse solve

```
n = 5000; e = ones(n,1);
A_dense = diag(-2*e) + diag(e(1:end-1),1) +
diag(e(1:end-1),-1);
A_sparse = spdiags([e -2*e e], -1:1, n, n);
b = ones(n,1);

t_dense = timeit(@() A_dense\b);
t_sparse = timeit(@() A_sparse\b);
[t_dense, t_sparse]
```

```
ans = 1x2
    0.0122    0.0001
```

Compare solve times on your machine (timeit reports seconds).

Create Sparse Matrices

Allocate space for $m \times n$ sparse matrix with v nonzeros:

```
S = spalloc(m; n; v)
```

Convert full matrix A to sparse matrix S:

```
S = sparse(A)
```

Create $m \times n$ sparse matrix with spare for v nonzeros from triplet (row,col,val):

```
S = sparse(row, col, val, m, n, v)
```

Create Sparse Matrices

Create matrix of 1s with sparsity structure defined by sparse matrix S:

$R = \text{spones}(S)$

Sparse identity matrix of size $m \times n$:

$I = \text{speye}(m, n)$

Create sparse uniformly distributed random matrix from sparsity structure of sparse matrix S:

$R = \text{sprand}(S)$

Create Sparse Matrices

Create sparse uniformly distributed random matrix of size $m \times n$ with approximately mnp nonzeros and condition number roughly κ (sum of rank 1 matrices):

```
R = sprand(m, n, rho, 1/kappa)
```

Create sparse normally distributed random matrix:

```
R = sprandn(S), R = sprandn(m, n, rho, 1/kappa)
```

Create sparse symmetric uniformly distributed random matrix:

```
R = sprandsym(S), R = sprandsym(n, rho, 1/kappa)
```

Import from sparse matrix external format:

```
spconvert
```

Create sparse matrices from diagonals:

```
spdiags
```

Sparse storage information

Determine if matrix is stored in sparse format:

`issparse(S)`

Number of nonzero matrix elements:

`nz = nnz(S)`

Amount of nonzeros allocated for nonzero matrix elements:

`nzmax(S)`

Extract nonzero matrix elements: If `(row, col, val)` is sparse triplet of `S`:

`val = nonzeros(S), [row,col,val] = find(S)`

Sparse and dense matrix functions

Convert sparse matrix to dense matrix:

```
A = full(S)
```

Apply function (described by function handle func) to nonzero elements of sparse matrix:

```
F = spfun(func, S)
```

Plot sparsity structure of matrix:

```
spy(S)
```

Sparse Matrix Reordering

By reordering the rows and columns of a matrix, it is possible to reduce the amount of fill-in that factorization creates, thereby reducing the time and storage cost of subsequent calculations.

Reordering Functions

`amd` Approximate minimum degree permutation

`colamd` Column approximate minimum degree permutation

`colperm` Sparse column permutation based on nonzero count

`dmp perm` Dulmage-Mendelsohn permutation/decomposition

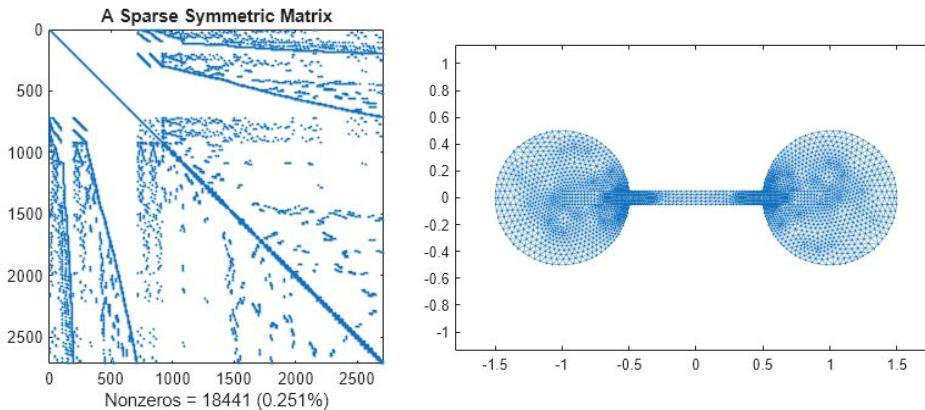
`randperm` Random permutation

`symamd` Symmetric approximate minimum degree permutation

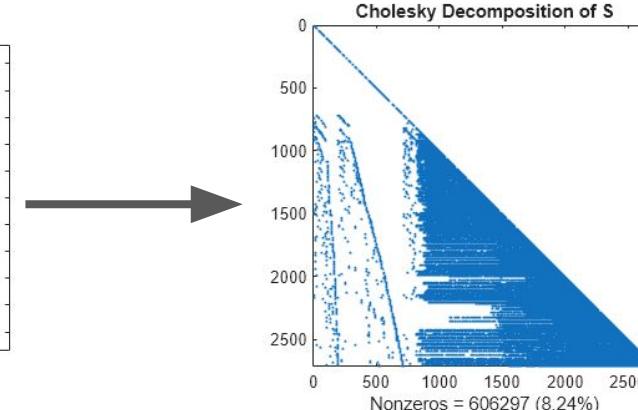
`symrcm` Sparse reverse Cuthill-McKee ordering

A Sparse Symmetric Matrix

<https://www.mathworks.com/help/matlab/math/sparse-matrix-reordering.html>



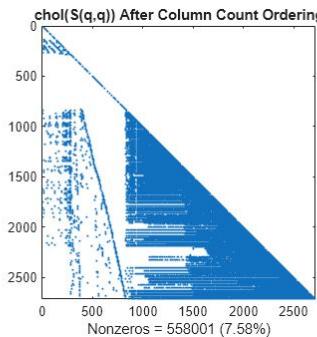
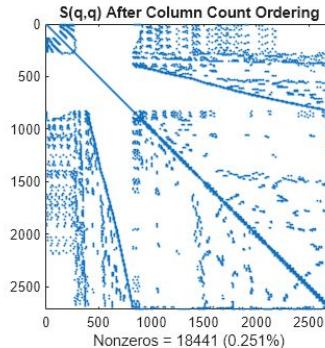
This `spy` plot shows a sparse symmetric positive definite matrix derived from a portion of the barbell matrix. This matrix describes connections in a graph that resembles a barbell.



This is the Cholesky factor L , where $S = L^*L'$. Notice that L contains many more nonzero elements than the unfactored S , because the computation of the Cholesky factorization creates fill-in nonzeros. These fill-in values slow down the algorithm and increase storage cost.

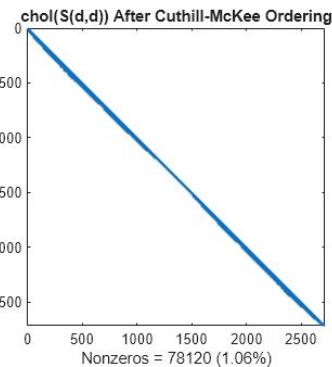
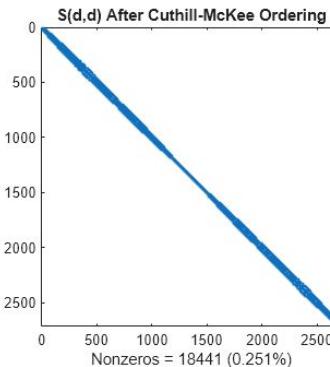
A Sparse Symmetric Matrix

By reordering the rows and columns of a matrix, it is possible to reduce the amount of fill-in that factorization creates, thereby reducing the time and storage cost of subsequent calculations.



The `colperm` command uses the column count reordering algorithm to move rows and columns with higher nonzero count towards the end of the matrix.

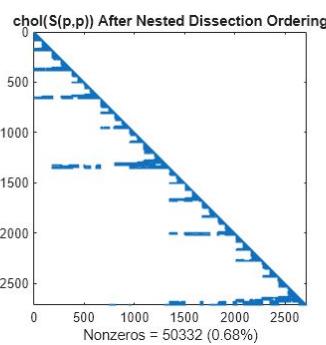
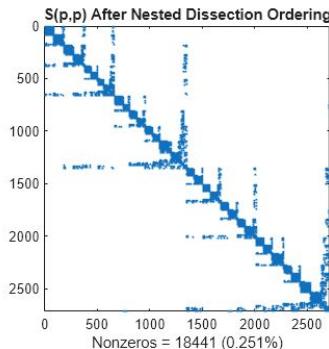
<https://www.mathworks.com/help/matlab/math/sparse-matrix-reordering.html>



The `symrcm` command uses the reverse Cuthill-McKee reordering algorithm to move all nonzero elements closer to the diagonal, reducing the bandwidth of the original matrix.

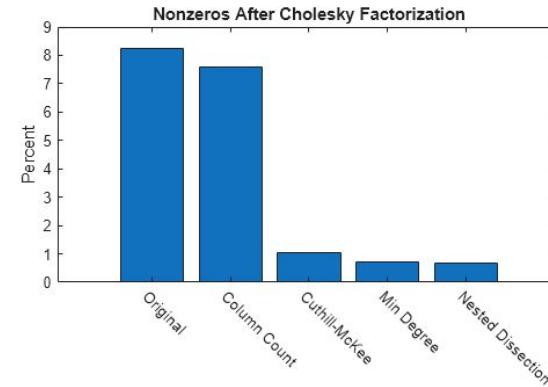
A Sparse Symmetric Matrix

By reordering the rows and columns of a matrix, it is possible to reduce the amount of fill-in that factorization creates, thereby reducing the time and storage cost of subsequent calculations.



The `dissect` function uses graph-theoretic techniques to produce fill-reducing orderings. The algorithm treats the matrix as the adjacency matrix of a graph, coarsens the graph by collapsing vertices and edges, reorders the smaller graph, and then uses refinement steps to uncoarsen the small graph and produce a reordering of the original graph. The result is a powerful algorithm that frequently produces the least amount of fill-in compared to the other reordering algorithms.

<https://www.mathworks.com/help/matlab/math/sparse-matrix-reordering.html>



This bar chart summarizes the effects of reordering the matrix before performing the Cholesky factorization. While the Cholesky factorization of the original matrix had about 8% of its elements as nonzeros, using `dissect` or `symamd` reduces that density to less than 1%.

Concrete example: reordering before factorization

```
S = sprandsym(2500, 0.002, 1e-2);
p1 = colperm(S);
p2 = amd(S);

R0 = chol(S);
R1 = chol(S(p1,p1));
R2 = chol(S(p2,p2));

[nnz(R0), nnz(R1), nnz(R2)]
```

```
ans = 1x3
      19658      7312      7280
```

Sparse Matrix Tips

Don't change sparsity structure (pre-allocate)

- Do not want to dynamically grows triplet
- Each component of triplet must be stored contiguously

Accessing values may be slow in sparse storage as location of row/columns is not predictable.

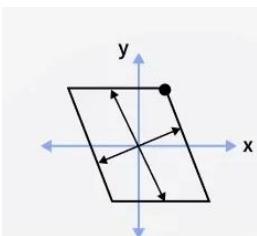
- If $S(i, j)$ is requested, must search through row, col to find i, j

Component-wise indexing to assign values is expensive

- Requires accessing into an array
- If $S(i, j)$ is previously zero, then $S(i, j) = c$ changes sparsity structure

Solving Linear Systems

Applied Math with MATLAB

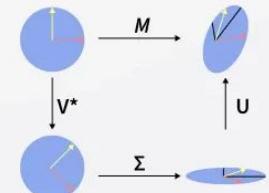


Linear Transformation

$$\begin{array}{c} \text{n} \times \text{n} \text{ Matrix} \\ \text{A} \end{array} \begin{array}{c} \xrightarrow{\quad} \\ \text{X} \end{array} = \begin{array}{c} \text{Eigenvalue} \\ \lambda \end{array} \begin{array}{c} \xrightarrow{\quad} \\ \text{X} \end{array}$$

Eigenvector Eigenvector

Eigenvalues and EigenVectors



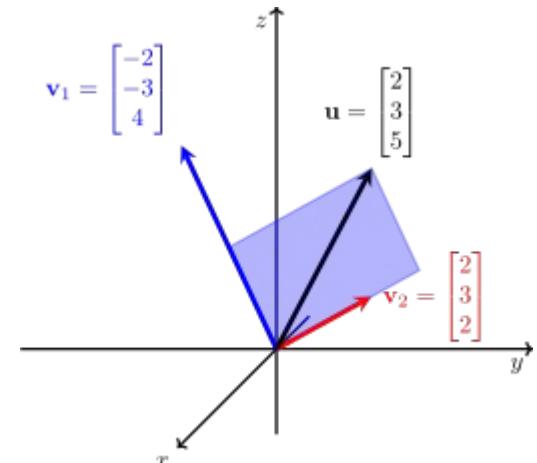
$$M = U \cdot \Sigma \cdot V^*$$

Singular Value Decomposition

Some basic concepts in linear algebra:

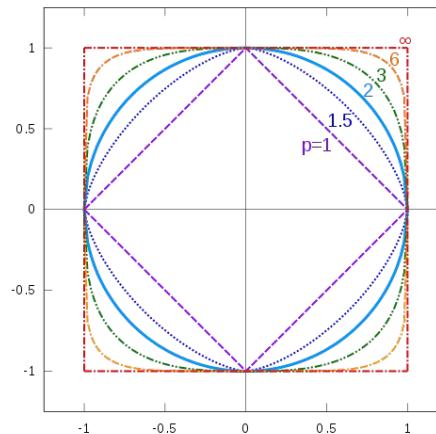
Rank

- the dimension of the vector space generated (or spanned) by its columns
- the maximal number of linearly independent columns
- the dimension of the vector space spanned by its rows



Norm

- 1-norm
- 2-norm
- Infinity-norm
- P-norms
- Frobenius norm



Linear system $Ax = b$ can be solved by factorizing the matrix A.

- Decompose A as $A = BC$, where B and C are matrices such that these two systems are easy to solve.
- Reduce the problem to solving $By = b$ and $Cx = y$.
- Examples of easy-to-solve matrices: diagonal, triangular, orthogonal
- For overdetermined system of equations, solve the linear least squares problem $\min 1/2 \|Ax - b\|_2^2$.

Hence, matrix decomposition!

LU Decomposition

$$A = LU$$

where A is non-singular, L is lower-triangular, and U is upper triangular.

Pivoting

- Gaussian elimination is unstable without pivoting.
- Partial pivoting: $PA = LU$
 - Permute the rows of A using P, such that the largest entry of the first column is at the top of that first column.
 - Apply Gaussian elimination without pivoting to PA.
- Complete pivoting: $PAQ = LU$
- Rook pivoting

Cholesky Factorization

$$A = R^*R = LL^*$$

where R is upper triangular and L is lower triangular ($*$: conjugate transpose).

A needs to be a Hermitian, positive-definite matrix.

- Cholesky Factorization is a variant of Gaussian elimination (LU) that operations on both left and right of the matrix simultaneously.
- Cholesky decomposition uses symmetric Gaussian elimination.
- Every Hermitian positive definite A has a unique Cholesky factorization.

Hermitian matrix $A = A^*$

Symmetric, positive definite (SPD) matrix

- A symmetric matrix A is SPD iff all its eigenvalues are positive → check by eigenvalue decomposition (expensive/difficult for large matrices)
- If a Cholesky decomposition can be successfully computed, the matrix is SPD → check by Cholesky factorization (best option)

QR Factorization

$$A = QR, \quad AE = QR$$

where Q is orthogonal ($QQ^T = I$), and R is upper triangular.

When is this useful?

- Pseudo-inverse
- Solution of least squares
- Solution of linear system of equations
- Extraction of orthogonal basis for column space of A

Full QR factorization

- Q: m x m, R: m x n

Economy QR (skinny QR) factorization

- Q: m x n, R: n x n
- Least-Squares: $\min \|Ax - b\|_2 = \min \|Rx - Q^T b\|_2$

Example: least-squares via QR / backslash

- Goal: solve $\min ||Ax - b||_2$ for an overdetermined system ($m > n$).
- Backslash uses a QR-based least-squares path for rectangular A.

```
% Fit y ≈ c1*x + c0
x = (0:0.1:1)';
A = [x ones(size(x))];
b = sin(2*pi*x);
c = A\b;
c

[Q,R] = qr(A,0);
c_qr = R\ (Q'*b);
c_qr
```

c =

2x1
-1.3989
0.6995

c_qr =

2x1
-1.3989
0.6995

Eigenvalue Decomposition (EVD)

$$A = XDX^{-1}$$

where D is a diagonal matrix with the eigenvalues of A on the diagonal and the columns of X contain the eigenvectors of A.

- Diagonalizable (EVD exists) vs. defective (EVD does not exist)
- All EVD algorithms must be iterative
- Eigenvalue Decomposition algorithm: reduce to upper Hessenberg form and iteratively transform upper Hessenberg to upper triangular

Quiz

```
A = gallery('lehmer', 4);
```

- Find the largest eigenvalue of the given matrix.
- Compute $\|A^3\|_\infty$ using EVD.
- Compute $\|e^A\|_1$ using EVD.

Solution

Quiz

```
A = gallery('lehmer', 4);
```

- Find the largest eigenvalue of the given matrix.
- Compute $\|A^3\|_\infty$ using EVD.
- Compute $\|e^A\|_1$ using EVD.

Solution

```
[V, D] = eig(A); max(D)
norm(V*D.^3/V, 'inf')
norm(V*exp(1).^D/V, 1)
```

Singular Value Decomposition (the most important matrix decomposition in numerical linear algebra)

$$A = U\Sigma V^T$$

where U and V are orthogonal and Σ is diagonal with real, positive entries.

Why do we care?

- Works for any matrix (square or rectangular, rank-deficient or not)
- Reveals the intrinsic rank and dimensionality of data
- Provides the best low-rank approximation (Eckart–Young theorem)
- Foundation of:
 - PCA and dimensionality reduction
 - Least-squares and ill-posed problems
 - Data compression, denoising, and regularization

Singular Value Decomposition (the most important matrix decomposition in numerical linear algebra)

$$A = U\Sigma V^T$$

where U and V are orthogonal and Σ is diagonal with real, positive entries.

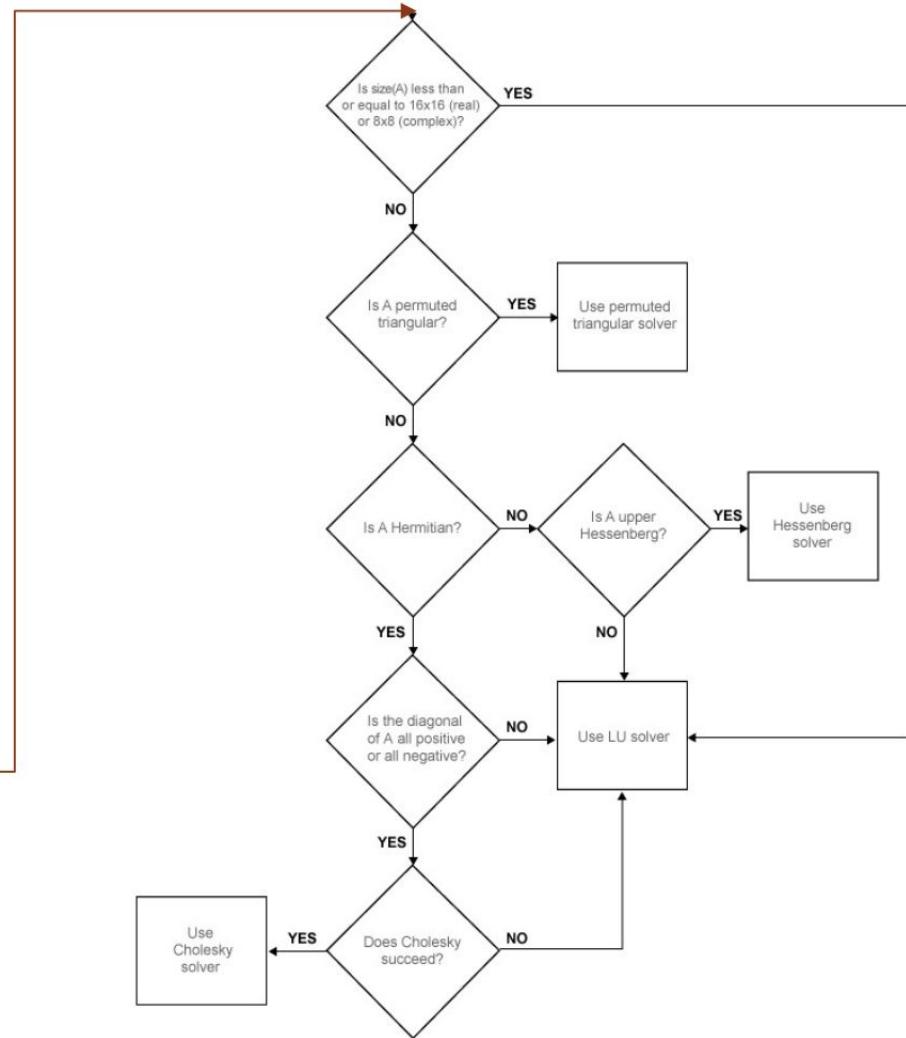
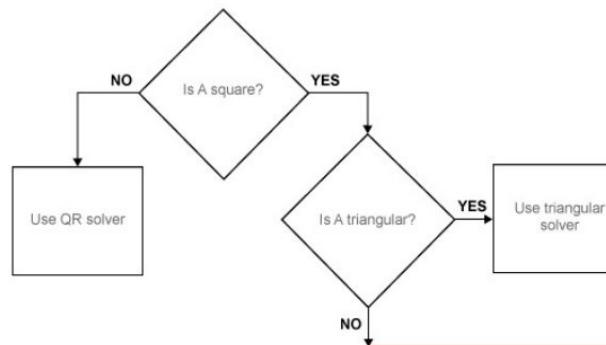
- SVD algorithm: Reduce A to bi-diagonal form and iteratively transform bi-diagonal to diagonal
- Full SVD: $A_{\{mxn\}} = U_{\{mxm\}} \Sigma_{\{mxn\}} V^T_{\{nxn\}}$
 - Use when you need complete orthonormal bases
 - Useful for theoretical analysis and proofs
 - More expensive in memory and computation
- Reduced SVD: $A_{\{mxn\}} = U_{\{mxr\}} \Sigma_{\{rxr\}} V^T_{\{rxn\}}$
 - $R = \text{rank}(A)$
 - Use for data analysis, PCA, compression
 - Faster, smaller, and almost always what you want in practice

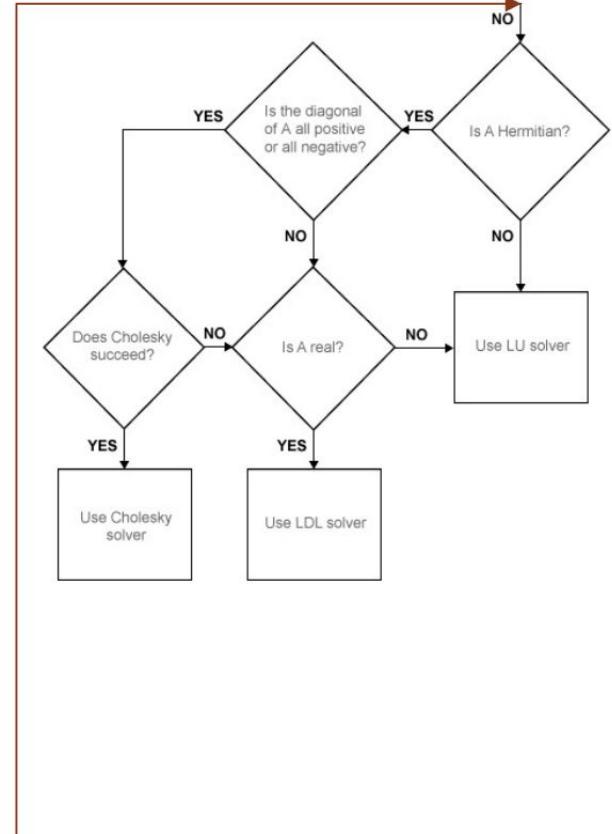
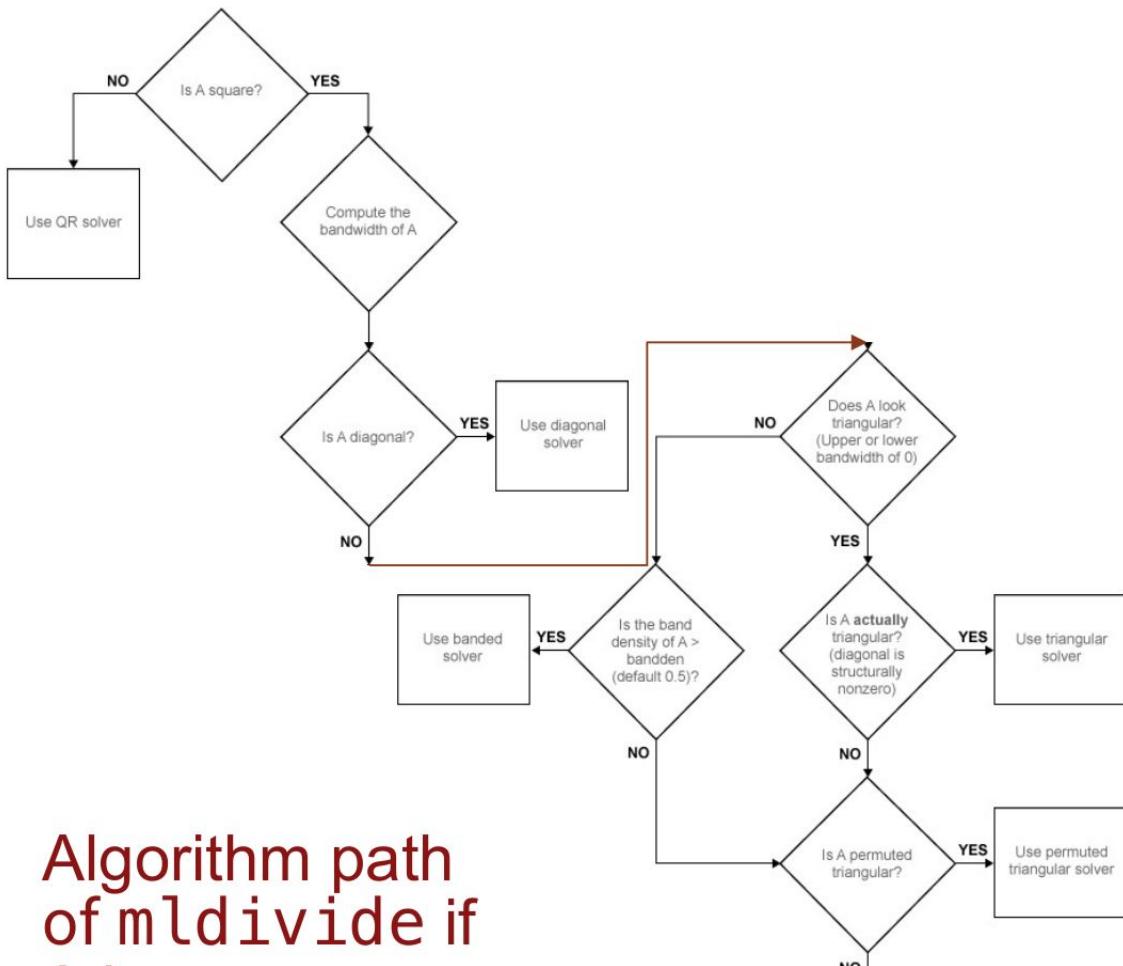
Direct Solvers for $Ax = b$: the Backslash

$$\begin{aligned}x &= A \backslash B \\x &= \text{mldivide}(A, B)\end{aligned}$$

- This solves the system of linear equations $A^*x = B$.
- The matrices A and B must have the same number of rows.
 - If A is a scalar, then $A \backslash B$ is equivalent to $A.\backslash B$.
 - If A is a square n-by-n matrix and B is a matrix with n rows, then $x = A \backslash B$ is a solution to the equation $A^*x = B$, if it exists.
 - If A is a rectangular m-by-n matrix with $m \approx n$, and B is a matrix with m rows, then $A \backslash B$ returns a least-squares solution to the system of equations $A^*x = B$.
- **Use backslash rather than $x = \text{inv}(A)^* b$**

Algorithm path of mldivide when A and B are full





Algorithm path of `mldivide` if A is sparse

Condition Number

A matrix is well-conditioned for condition number k close to 1;
ill-conditioned for condition number k large.

- `cond`: returns 2-norm condition number
- `condest`: lower bound for 1-norm condition number
- `rcond`: LAPACK estimate of inverse of 1-norm condition number

In $Ax = b$, if the condition number is large, even a small error in b may cause a large error in x .

Example: condition number sensitivity

- Small relative perturbation in b can cause larger relative change in x when $\text{cond}(A)$ is large.

```
A = hilb(10
b = ones(10,1)
x = A\b
condA = cond(A)
db = 1e-8 * randn(size(b))
x2 = A\b + db
rel_b = norm(db)/norm(b)
rel_x = norm(x2 - x)/norm(x)
fprintf('Condition number of A: %.2e\n',
condA); fprintf('Relative perturbation in
b: %.2e\n', rel_b)
fprintf('Relative change in solution:
%.2e\n', rel_x)
fprintf('Amplification factor (~cond):
%.2e\n', rel_x / rel_b)
```

```
Condition number of A: 1.60e+13
Relative perturbation in b: 8.55e-09
Relative change in solution: 4.55e-03
Amplification factor (~cond): 5.33e+05
```

Iterative Solvers

- **Goal:** Solve large linear systems $Ax=b$ approximately by iteratively improving the solution, without forming or factorizing A .
- **Who uses them:** Scientists and engineers in numerical PDEs, machine learning, optimization, physics simulations, and large-scale data analysis.
- **Why they matter:** They scale to massive sparse problems, use far less memory than direct solvers, and often converge quickly when combined with good preconditioners.

Iterative Solvers

Preconditioning

- Preconditioning replaces the original problem ($Ax = b$) with a different problems with the same (or similar) solution.
 - Left preconditioning: $L^{-1}Ax = L^{-1}b$
 - Right preconditioning: $y = Rx$, $AR^{-1}y = b$
 - Left and right preconditioning: $L^{-1}AR^{-1}y = L^{-1}b$
- Preconditioner M for A ideally provides a cheap approximation to A^{-1} , intended to drive condition number toward 1.
- Typical preconditioners:
 - Jacobi: $M = \text{diag}(\text{diag}(A))$
 - Incomplete factorizations: LU, Cholesky (control for level of fill-in)

Common Iterative Solvers

Linear system of equations $Ax = b$

- Symmetric Positive Definite matrix: Conjugate Gradients (CG)
- Symmetric matrix: Symmetric LQ Method (SYMMLQ), Minimum-Residual (MINRES)
- General, Unsymmetric matrix: Biconjugate Gradients (BiCG), Biconjugate Gradients Stabilized (BiCGstab), Conjugate Gradients Squared (CGS), Generalized Minimum-Residual (GMRES)

Linear least-squares $\min \|Ax - b\|_2$

- Least-Squares Minimum-Residual (LSMR)
- Least-Squares QR (LSQR)

Example: preconditioned CG (pcg) in MATLAB

- Constructs a sparse SPD matrix from a 2D Laplacian, then solves $A x = b$ using the preconditioned conjugate gradient method. The incomplete Cholesky factor serves as a preconditioner to significantly accelerate convergence compared to unpreconditioned CG.

```
n = 1000
m = round(sqrt(n)) + 2
A = delsq(numgrid('S', m))
b = ones(size(A,1),1)
L = ichol(A, struct('diagcomp',1e-2))
[x,flag,relres,iter,resvec] = pcg(A,b,1e-8,200,L,L')
fprintf('size(A)=%d, condest(A)=%.2e\n', size(A,1),
condest(A))
fprintf('pcg: flag=%d, iter=%d, relres=%.2e\n', flag, iter,
relres)
```

```
size(A)=1024, condest(A)=6.40e+02
pcg: flag=0, iter=30, relres=3.82e-09
```

MATLAB's built-in iterative solvers for $Ax = b$, $A \in \mathbb{R}^{m \times m}$

```
[x,flag,relres,iter,resvec] =  
solver(A,b,restart,tol,maxit,M1,M2,x0)
```

Inputs (only A, b required):

- A – full or sparse (recommended) square matrix or function handle returning Av for $v \in \mathbb{R}^m$
- b – m vector
- restart – restart frequency (GMRES)
- tol – relative convergence tolerance
- maxit – maximum number of iterations
- $M1, M2$ – full or sparse (recommended) preconditioner matrix or function handle returning $M2^{-1}M1^{-1}v$ for any $v \in \mathbb{R}^m$
- $x0$ – initial guess of solution

Outputs:

```
[x,flag,relres,iter,resvec] =  
solver(A,b,restart,tol,maxit,M1,M2,x0)
```

- x – attempted solution to $Ax = b$
- flag – convergence flag
- relres – relative residual $\|b - Ax\|/\|b\|$ at convergence
- iter – number of iterations (inner and outer iterations for certain algorithms)
- resvec – vector of residual norms at each iteration $\|b - Ax\|$, including preconditioners if used ($\|M^{-1}(b - Ax)\|$).

Symbolic Math

Applied Math with MATLAB

MATLAB speaks ...

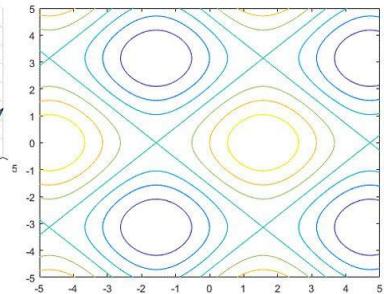
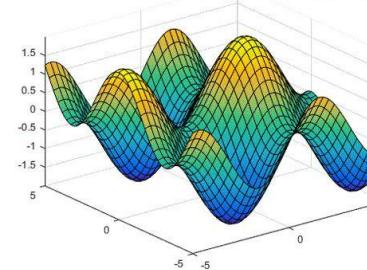
$$f(x) = x^4 - 2 * x^3 + 6 * x^2 - 2 * x + 10$$

$$A * x = b \text{ where } A = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \text{ and } x = [x_1, x_2]$$

$$f(x, y) = \sin(x) + \cos(y)$$

$$\sum (x-a)^n \frac{f^{(n)}(a)}{n!} \text{ for } f(x) = \frac{\sin(x)}{x}$$

$$\lim_{x \rightarrow \frac{\pi}{2}^+} \tan(x) \neq \lim_{x \rightarrow \frac{\pi}{2}^-} \tan(x).$$



Symbolic Math Toolbox provides functions for solving, plotting, and manipulating symbolic math equations.

- analytically perform differentiation, integration, simplification,
- transforms, and equation solving
- perform dimensional computations and convert between units
- display computation results in mathematical typeset
- convert work to HTML, Word, LaTex, or PDF documents

Operations and Commands

Symbolic arithmetic operations

- ceil, cong, cumprod, cumsum, fix, floor, frac, imag,
- minus, mod, plus, quorem, real, round

Symbolic relational operations

- eq, ge, gt, le, lt, ne, isequaln

Symbolic logical operations

- and, not, or, xor, all, any, isequaln, isfinite, isinf,
- isnan, logical

Equation Solving

- `finvrse` Functional inverse
- `linsolve` Solve linear system of equations
- `poles` Poles of expression/function
- `solve` Equation/System of equations solver
- `dsolve` ODE solver

Formula Manipulation and Simplification

- `simplify` Algebraic simplification
- `simplifyFraction` Symbolic simplification of fractions
- `subexpr` Rewrite symbolic expression in terms of common subexpression
- `subs` Symbolic substitution

Calculus

- `diff` Differentiate symbolic
- `int` Definite and indefinite integrals
- `rsums` Riemann sums
- `curl` Curl of vector field
- `divergence` Divergence of vector field
- `gradient` Gradient vector of scalar function
- `hessian` Hessian matrix of scalar function
- `jacobian` Jacobian matrix
- `laplacian` Laplacian of scalar function
- `potential` Potential of vector field
- `vectorPotential` Vector potential of vector field
- `taylor` Taylor series expansion
- `limit` Compute limit of symbolic expression
- `fourier` Fourier transform
- `ifourier` Inverse Fourier transform
- `ilaplace` Inverse Laplace transform
- `iztrans` Inverse Z-transform
- `laplace` Laplace transform
- `ztrans` Z-transform

$$\int_0^{\infty} e^{-ax^2} \cos(bx) dx = \frac{\sqrt{\pi}}{2\sqrt{a}} e^{-b^2/(4a)}, \quad a > 0.$$

```

syms x a b positive real
I = int(exp(-a*x^2)*cos(b*x), x, 0, inf);
I_simplified = simplify(I)
disp('Integral result:')
pretty(I_simplified)

a_val = 2
b_val = 3;
I_num_sym = double(subs(I_simplified, [a b],
[a_val b_val]))
I_num_quad = integral(@(t)
exp(-a_val*t.^2).*cos(b_val*t), 0, Inf)

fprintf('Symbolic: %.15f\n', I_num_sym)
fprintf('Numeric : %.15f\n', I_num_quad)
fprintf('Abs diff: %.3e\n', abs(I_num_sym -
I_num_quad))

```

Integral result:

$$\frac{\sqrt{\pi} \exp\left(-\frac{b^2}{4a}\right)}{2\sqrt{a}}$$

Symbolic: 0.203445763527289
 Numeric : 0.203445763527288
 Abs diff: 1.027e-15

Linear Algebra

Most matrix operations available for numeric arrays also available for symbolic matrices.

- `adjoint` Adjoint of symbolic square matrix
- `expm` Matrix exponential
- `sqrtm` Matrix square root
- `cond` Condition number of symbolic matrix
- `det` Compute determinant of symbolic matrix
- `norm` Norm of matrix or vector
- `colspace` Column space of matrix
- `null` Form basis for null space of matrix

Linear Algebra

- `rank` Compute rank of symbolic matrix
- `rref` Compute reduced row echelon form
- `eig` Symbolic eigenvalue decomposition
- `jordan` Jordan form of symbolic matrix
- `chol` Symbolic Cholesky decomposition
- `lu` Symbolic LU decomposition
- `qr` Symbolic QR decomposition
- `svd` Symbolic singular value decomposition
- `inv` Compute symbolic matrix inverse
- `linsolve` Solve linear system of equations

Assumptions

- `assume` Set assumption on symbolic object
- `assumeAlso` Add assumption on symbolic object
- `assumptions` Show assumptions set on symbolic variable

Polynomials

- `charpoly` Characteristic polynomial of matrix
- `coeffs` Coefficients of polynomial
- `minpoly` Minimal polynomial of matrix
- `poly2sm` Symbolic polynomial from coefficients
- `sym2poly` Symbolic polynomial to numeric

$$\int_0^\infty \frac{x^{s-1}}{1+x} dx = \pi \csc(\pi s), \quad 0 < s < 1.$$

```

syms x s real
assumeAlso(s > 0 & s < 1)

I = int(x^(s-1)/(1+x), x, 0, inf);
I_simplified = simplify(I);

disp('Integral result:')
pretty(I_simplified)

s_val = 0.37;
I_num_sym = double(subs(I_simplified, s, s_val));
I_num_quad = integral(@(t) t.^(s_val-1)./(1+t),
0, Inf);

fprintf('Symbolic: %.15f\n', I_num_sym);
fprintf('Numeric : %.15f\n', I_num_quad);
fprintf('Abs diff: %.3e\n', abs(I_num_sym -
I_num_quad));

```

Integral result:

$$\frac{\pi}{\sin(\pi s)}$$

 Symbolic: 3.423129195615263
 Numeric : 3.423128249625757
 Abs diff: 9.460e-07

Mathematical Functions

- `log`, `log10`, `log2` Logarithmic functions
- `sin`, `cos`, `tan`, etc Trigonometric functions
- `sinh`, `cosh` `tanh`, etc Hyperbolic functions

Precision Control

- `digits` Variable-precision accuracy
- `double` Convert symbolic expression to MATLAB double
- `vpa` Variable precision arithmetic

Code generation

- `ccode` C code representation of symbolic expression
- `fortran` Fortran representation of symbolic expression
- `latex` LATEX representation of symbolic expression
- `matlabFunction` Convert symbolic expression to function handle or file
- `texlabel` TeX representation of symbolic expression